

SALES PREDICTION USING PYTHON

Objective

The goal is to predict product sales based on advertising expenditure across three platforms: TV, Radio, and Newspaper. By understanding the relationship between advertising efforts and sales, businesses can make informed decisions on how to allocate their advertising budgets effectively.

Dataset Description

The dataset contains the following columns:

- **TV:** Advertising expenditure on TV campaigns (in thousands of dollars).
- **Radio:** Advertising expenditure on Radio campaigns (in thousands of dollars).
- **Newspaper:** Advertising expenditure on Newspaper campaigns
- **Sales:** Total product sales (in thousands of units).

Key Steps

- **Import Libraries:** libraries for data handling, visualization, and machine learning.
- **Load Data:** Load the dataset, check for missing values, and summarize statistics.
- **EDA:** Visualize relationships (scatter plots, heatmaps) and analyze correlations.
- **Feature and Target Setup:** Define independent variables (X) and the target (y).
- **Split Data:** Divide the dataset into training and testing sets using `train_test_split()`.
- **Train Model:** Fit a Linear Regression model to the training data.
- **Model Parameters:** Output the model's intercept and coefficients.
- **Make Predictions:** Predict sales on the test dataset.
- **Evaluate Model:** Calculate R-squared and RMSE to assess model performance.
- **Visualize Results:** Compare actual vs predicted sales using a scatter plot.
- **Conclude:** Summarize findings and recommend strategies based on insights.

Packages

- **pandas:** Used to load and manipulate data (e.g., reading CSV files, inspecting data).
- **numpy:** Used for numerical operations (e.g., calculating RMSE).
- **matplotlib.pyplot:** Used for plotting graphs (e.g., scatter plots, line plots).
- **seaborn:** Built on matplotlib, used for creating more advanced and aesthetically pleasing visualizations (e.g., pairplots, heatmaps).
- **sklearn.model_selection.train_test_split:** Splits the dataset into training and testing set
- **sklearn.linear_model.LinearRegression:** Used to create and train the linear regression model.
- **sklearn.metrics.mean_squared_error & r2_score:** Used to evaluate the model's performance (e.g., RMSE, R^2 score).

Step 1: Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

Step 2: Loading the Dataset

```
data = pd.read_csv('sales_pred.csv')
# Data Preprocessing and Exploration
print(data.head()) # Display the first 5 rows of the dataset
print(data.isnull().sum()) # Check for missing values in each column
print(data.describe()) # Get a summary of the numerical columns (mean, min, max, etc.)
```

Output:

First 5 rows of the dataset:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

Missing values in each column:

TV	0
Radio	0
Newspaper	0
Sales	0

dtype: int64


Statistical description of the dataset:

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

Step 3: Visualize Relationships


Pairplot to see relationships between features and target

```
Visualize relationships between features (TV, Radio, Newspaper) and Sales using a pairplot
sns.pairplot(data, x_vars=['TV', 'Radio', 'Newspaper'], y_vars='Sales', height=4,
aspect=1, kind='scatter')
plt.show() # Display the pairplot
```



Correlation Heatmap

```
Visualize the correlation matrix using a heatmap
sns.heatmap(data.corr(), annot=True, cmap="coolwarm") # Display correlation
matrix with annotations
plt.show()
```



Step 4: Splitting Data into Features and Target Variable

```
Split the dataset into features (X) and target variable (y)
X = data[['TV', 'Radio', 'Newspaper']] # Independent variables (features)
y = data['Sales'] # Dependent variable (target: Sales)

# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

Step 5: Training the Linear Regression Model

```
Initialize and train the Linear Regression model
model = LinearRegression() # Create a linear regression model
model.fit(X_train, y_train) # Fit the model to the training data

# Output the intercept and coefficients of the trained model
print(f"Intercept: {model.intercept_}") # Display the intercept (constant term)
print(f"Coefficients: {model.coef_}")
```

Output

Intercept: 4.714126402214127

Coefficients: [0.05450927 0.10094536 0.00433665]

Step 6: Making Predictions

```
y_pred = model.predict(X_test) # Predict the sales on the test set
```

Step 7: Evaluating the Model

```
Evaluate the model using R-squared and RMSE (Root Mean Squared Error)
r2 = r2_score(y_test, y_pred) # Calculate R-squared score (goodness of fit)
rmse = np.sqrt(mean_squared_error(y_test, y_pred)) # Calculate RMSE (lower is
better)

print(f"R-squared: {r2:.2f}") # Display the R-squared value
print(f"RMSE: {rmse:.2f}") # Display the RMSE value
```

Output


R-squared: 0.91

RMSE: 1.71

Step 8: Visualizing Predictions

```
Visualize the comparison of actual vs predicted sales using a scatter plot
plt.scatter(y_test, y_pred, alpha=0.8) # Scatter plot for actual vs predicted
```

```
plt.xlabel("Actual Sales") # Label for the x-axis
plt.ylabel("Predicted Sales") # Label for the y-axis
plt.title("Actual vs. Predicted Sales") # Title of the plot
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color="red") #
Line of perfect prediction
plt.show()
```



Conclusion

The analysis reveals that advertising on TV and Radio had the most significant impact on sales, as indicated by both the coefficients and the correlation heatmap. Newspaper advertising, however, showed a relatively lesser influence, suggesting an opportunity for optimizing budget allocation. The model's performance is promising, with a high R-squared value, indicating that the model explains a significant portion of the variance in sales. Additionally, the RMSE highlights the average deviation between the predicted and actual sales. By leveraging this linear regression model, businesses can predict sales based on their advertising investments, enabling data-driven decisions for more efficient budget allocation and maximizing return on investment (ROI).

